Who Carries the Cartel?
LLM Price Collusion Survives Agent Replacement Through the Inherited Price History

Supratik Bhowal
Department of Computer Science (AIML)
IEM Kolkata, School of UEMK, Kolkata, India
supratikbhowal23@gmail.com

Abstract

When one agent of an LLM pricing cartel is replaced, is collusion carried by the agent or by the price history the entrant inherits? In a pre-registered experiment in a logit Bertrand duopoly, we copy each Qwen2.5-7B cartel at round 100 into paired branches that differ only in the entrant's model, visible history and notes. In six independent runs (16–31 sessions each), a new instance of the same model that inherits the history keeps the cartel (profit index within -0.025 to $+0.213$ of no replacement); without the history it loses 0.47–0.67, mostly by pricing above the monopoly price. With the history, a Mistral-7B entrant drives prices to the competitive level, while an OLMo-2-7B entrant keeps them, repeating its last price in 96% of rounds. A forced-deviation audit calibrated against a myopic control finds almost no punishment. The cartel lives in the transcript, and it is anchored, not enforced.

1 Introduction

Large language models (LLMs) used as pricing agents reach supra-competitive prices without communication and without instructions to collude [Fish et al., 2024], as Q-learners do in the same market [Calvano et al., 2020]. A natural remedy, and an event that happens whenever a firm changes its software vendor or model version, is to *replace* one agent. Whether that ends the cartel depends on where the cartel is stored. A prompted LLM agent has no state outside its prompt: a window of past prices and the notes it wrote to itself. An entrant with the same model, the same history and the same notes *is* the old agent. The question therefore splits exactly: is the cartel carried by the *model*, by the inherited *price history*, or by the inherited *notes*?

We answer it with paired branches. Two Qwen2.5-7B agents play 100 rounds; every session that reaches a cartel-like level is copied into eight branches that differ only in what occupies seat 1 from round 100 (Table 1). Sampling seeds do not depend on the branch, so branches of one session differ only through the swap. A forced-deviation audit [Tran and Yu, 2026] then tests whether the surviving prices are enforced or anchored.

Our main contributions are:

1. **The price history carries the cartel.** In six independent runs, a new instance of the same model that inherits the history keeps the cartel, and without the history it does not (H1, H2, H3a supported in 6 of 6). No swap notice is needed (H7), and the predecessor's notes alone keep part of the cartel (H9, primary test).
2. **Removing the history does not restore competition.** The blank entrant lowers profit by pricing *above* the monopoly price, which a profit index alone misreads.

3. **The entrant's pricing style decides the outcome.** With the history, a Mistral-7B entrant drives prices to the competitive level (4 of 4 runs); an OLMo-2-7B entrant holds its last price in 96% of rounds and keeps the cartel (2 of 2 runs).
4. **No enforcement after calibration.** Compared with a myopic control that cannot punish, almost no audit cell shows punishment. We contribute this calibration and request deduplication for exactly paired LLM branches.

2 Related Work

2.1 LLM Price Collusion and Enforcement

[Calvano et al., 2020] show that Q-learners in a repeated logit Bertrand oligopoly learn supra-competitive prices sustained by finite punishments. [Fish et al., 2024] move this market to LLM agents, which reach supra-competitive prices quickly, with a level that depends on prompt wording and with text that voices price-war concerns. Both keep the same agents for the whole game; our market and calibration follow theirs, and our treatment is replacement. Because a cartel is defined by enforcement, [Tran and Yu, 2026] audit it with paired forced deviations under shared seeds and find that prompted open models rarely punish: their high prices come from rules that ignore the rival or anchor on the price scale. They also show that a pure best responder follows a price cut by 0.29–0.44 per unit, the benchmark we subtract.

2.2 Partners, Replacement and Heterogeneity

[Arslan et al., 2026] show in a pre-registered experiment that keeping the same partner raises learned prices; in their prompted extension, Mistral-7B and OLMo-2-7B price high in both arms. [Gao et al., 2026] swap agents in cooperative LLM teams with a placebo and cleared notes, a design we transfer to collusion. [Keppo et al., 2026] find that heterogeneity in patience, data access and algorithm weakens LLM collusion, while model size does not. [Lee, 2026] find that collusion is a property of the model, and that GPT markets earn no collusive rent although their prices drift above the monopoly level. None of these studies replaces an agent inside an established LLM cartel.

3 Methods

3.1 Market and Agents

Two firms face logit demand $q_i = e^{(a-p_i)/\mu} / \left(\sum_j e^{(a-p_j)/\mu} + e^{a_0/\mu} \right)$ with profit $\pi_i = (p_i - c)q_i$, $a = 2$, $a_0 = 0$, $\mu = 0.25$, $c = 1$ [Calvano et al., 2020]. The Nash price is $p^N = 1.473$ and the joint-monopoly price is $p^M = 1.925$. We report the profit index $\Delta = (\bar{\pi} - \pi^N)/(\pi^M - \pi^N)$ (pre-registered outcome) and the price index $P = (\bar{p} - p^N)/(p^M - p^N)$ (pre-registered secondary outcome), both means over the two firms, 0 at Nash and 1 at monopoly. A price above p^M gives a high P but a low or negative Δ.

Each firm is a prompted LLM served with vLLM 0.8.5 (temperature 0.7, at most 160 output tokens). Each round it sees the market rules in plain words, the goal "maximize your firm's total profit over all periods", its last 30 rounds, and its previous notes, and replies with JSON holding new notes and a price (Appendix D). The prompt never mentions competition, collusion, coordination or punishment. Incumbents are Qwen2.5-7B-Instruct; the other families are Mistral-7B-Instruct-v0.3 and OLMo-2-1124-7B-Instruct.

3.2 Design and Pairing

In phase 1, two incumbents play 100 rounds (30 or 40 sessions per run). A session is kept when its mean profit index over rounds 50–99 exceeds 0.3 (pre-registered). In phase 2, every kept session is copied into the branches of Table 1 and plays 100 more rounds. In swap branches the incumbent is told "From period 100 on, the other firm uses a new pricing system"; `placebo` gives the notice without a swap.

Table 1: Phase-2 branches: what occupies seat 1 from round 100, the part of the price history it sees, and its notes.

Branch	Seat-1 agent	History seen	Notes
`control`	no swap	all	kept
`placebo`	no swap; incumbent gets the notice	all	kept
`same_inherit`	Qwen2.5-7B, new instance	last 30 rounds	cleared
`same_blank`	Qwen2.5-7B, new instance	from round 100	cleared
`same_notes`	Qwen2.5-7B, new instance	from round 100	predecessor's
`other_inherit`	other family	last 30 rounds	cleared
`other_blank`	other family	from round 100	cleared
`br_bot`	best response to rival's last price	last rival price	–

Each request carries the seed $h(\text{experiment seed}, \text{session}, \text{seat}, \text{round}, \text{attempt})$, where h is a hash that never takes the branch, so branches that show an agent the same prompt sample the same reply and all comparisons are paired per session. GPU numerics depend on the batch, so this pairing is inexact in large batches (Section 4.1); from run 9 on, identical requests in a round are generated once and the reply is copied (request de-duplication). The A/A check runs without it.

3.3 Enforcement Audit and Its Calibration

At round 200, every session is copied twice per deviator seat; in the "cut" copy the deviator's price is forced 10% below its last price for one round, and both copies play 6 more rounds with the same seeds. The excess response is $e = r - r^{BR}$, where r is the rival's price cut one round later per unit of the deviator's cut (base copy minus cut copy) and r^{BR} is the same for a best responder [Tran and Yu, 2026]. Three prompt controls (10 sessions per variant in runs 4–8, 30 in runs 9–12) validate the audit: the default prompt; a *myopic* prompt ("profit in future periods does not count"), which cannot punish; and an instructed *trigger* rule (charge the monopoly price, drop to the Nash price after a cut), which must be detected. Since LLM agents follow cuts even when they cannot punish, test H5c counts a cell as punishment only if the 95% two-sample bootstrap interval of $\hat{e}_{\text{cell}} - \hat{e}_{\text{myopic}}$ is above zero.

3.4 Hypotheses and Statistics

Hypotheses, inclusion rule and metrics were fixed in a pre-registration committed to the project repository before the first main run (22 Oct 2026, 18:38 UTC); later hypotheses were added before the runs that test them. Differences are paired per kept session, in the post window (rounds 150–199), with 10,000 bootstrap resamples over sessions. A hypothesis is *supported* if the whole 95% interval lies on its stated side of the threshold, *not supported* if it lies on the other side, and *inconclusive* otherwise (Table 2). All runs are reported (Appendix Table 6): six with a Qwen2.5-7B

incumbent (runs 4, 6–10), two whose control cartel collapsed (runs 5 and 12), and one that failed (run 11).

Table 2: Pre-registered hypotheses. `si`, `sb`, `oi`: `same_inherit`, `same_blank`, `other_inherit`.

Test	Claim	Meaning
H1	`si` − control > -0.10	history carries the cartel
H2	`sb` − control < -0.10	a blank entrant breaks it
H3a / H3b	`si` − `sb` > 0 / `oi` − `other_blank` > 0	history matters for each family
H4	$\lvert$`si` − `sb`$\rvert - \lvert$`si` − `oi`$\rvert > 0$	history matters more than model
H5 / H5c	no LLM cell with e above 0 / above the myopic control	no enforcement
H6 / H8	`oi` − control < -0.10 / its opposite for OLMo-2 (run 10)	other family breaks / keeps it
H7	H1 without the swap notice (run 8)	no notice needed
H9	`same_notes` − `sb` > 0	notes alone carry part of it

4 Results

4.1 Controls and Data Quality

A/A determinism was 1.0 in every run except run 9 (0.8125). Parse failures were at most 1.0% per phase-1 session and 0.3% per phase-2 branch in every run with a Qwen2.5-7B incumbent. The `br_bot` negative control lowered the price index by 0.715–0.949 against control, to the Nash level. The calibrated trigger control was detected in runs 4, 6, 7, 8 and 9 (trigger minus myopic 1.316, 0.638, 0.760, 1.044, 0.361, all intervals above zero) but not in run 10 (0.269 [−0.061, 0.598]), so by the pre-registered rule the audit of run 10 is not interpreted. Trigger agents punish fully or not at all (exploratory): one round after the cut, the trigger agent's price was 0.45 lower in the cut copy in 70 of 138 events, unchanged in 67 and 0.45 higher in 1. A drop of at least 0.2 beyond the best response, our punishment criterion, occurred in 4–7 of the 9–10 sessions per seat in runs 4–8, and in 11 of 29 and 11 of 30 in run 9. Without de-duplication, pairing was weaker than the A/A check suggests: in run 8, where `placebo` has the same design as control, only 47.9% and 46.0% (seats 0 and 1) of post-swap decisions were identical, because one different sample changes every later prompt. The mean effect was unaffected (−0.001 [−0.022, 0.015]), so runs 4–8 are unbiased, with wider intervals.

4.2 The Price History Carries the Cartel

In all six runs with a Qwen2.5-7B incumbent, a new instance of the same model that inherits the price history stays at the control level, and the same model without the history falls far below it (Table 3, Figure 1). Pooling the four runs that share the design of these three branches (runs 4, 6, 7, 9; $n = 92$; exploratory) gives H1 = 0.053 [−0.011, 0.148], H2 = −0.592 [−0.751, −0.432] and H3a = 0.645 [0.503, 0.793]. Without the swap notice (run 8), H1 is 0.002 [−0.014, 0.016] (H7 supported). In run 9 the H1 interval also lies below zero (−0.025 [−0.051, −0.001]): with the history, the cartel continues slightly but measurably lower. In run 7, `same_inherit` is above control because that control cartel weakened on its own (0.930 before the swap, 0.687 after).

Table 3: Main result, every run with a Qwen2.5-7B incumbent: paired differences in the post window, estimate [95% CI]. †Run 10 values come from its verbatim summary in the run log.

Run	n	H1: same_inherit − control	H2: same_blank − control	H3a: same_inherit − same_blank	Price index, same_blank − control
4	16	−0.003 [−0.065, 0.064]	−0.670 [−0.998, −0.360]	0.667 [0.368, 0.989]	+0.492 [0.162, 0.834]
6	22	0.016 [−0.077, 0.122]	−0.474 [−0.782, −0.202]	0.489 [0.247, 0.771]	+0.582 [0.292, 0.888]
7	25	0.213 [0.006, 0.529]	−0.624 [−1.021, −0.215]	0.836 [0.553, 1.164]	+0.406 [−0.055, 0.878]
8	24	0.002 [−0.014, 0.016]	−0.504 [−0.850, −0.208]	0.505 [0.209, 0.853]	+0.337 [0.036, 0.663]
9	29	−0.025 [−0.051, −0.001]	−0.611 [−0.856, −0.380]	0.587 [0.349, 0.838]	+0.481 [0.223, 0.746]
10†	31	−0.014 [−0.038, 0.008]	−0.480 [−0.689, −0.292]	0.467 [0.279, 0.676]	+0.420 [0.163, 0.686]
Verdicts		supported in 6 of 6	supported in 6 of 6	supported in 6 of 6	CI above 0 in 5 of 6

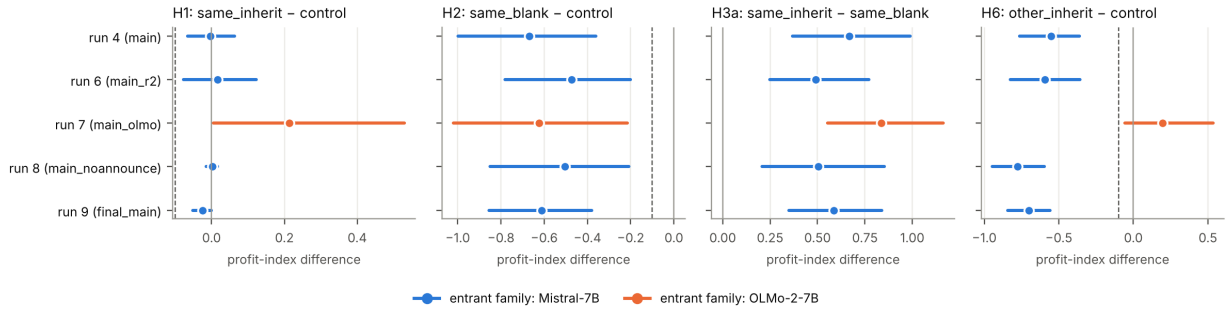


Figure 1: Pre-registered tests per run (estimate, 95% CI; dashed: the −0.10 thresholds). H1, H2 and H3a hold in every run; H6 depends on the entrant family (blue: Mistral-7B, orange: OLMo-2-7B). Run 10 is not drawn because its run folder was not available to the figure script.

4.3 Removing the History Does Not Restore Competition

The blank entrant's profit loss is not a return to competition: its price index is above control in all six runs, with the interval above zero in five (Table 3). Without a history, the new agent starts high and stays above the monopoly price, and both firms earn less. Figure 2 shows run 9: `same_blank` and `same_notes` sit between the Nash and monopoly profit in panel (a), while panel (b) shows prices above the monopoly level throughout. Only `br_bot` and the Mistral entrant with the history reach the Nash level.

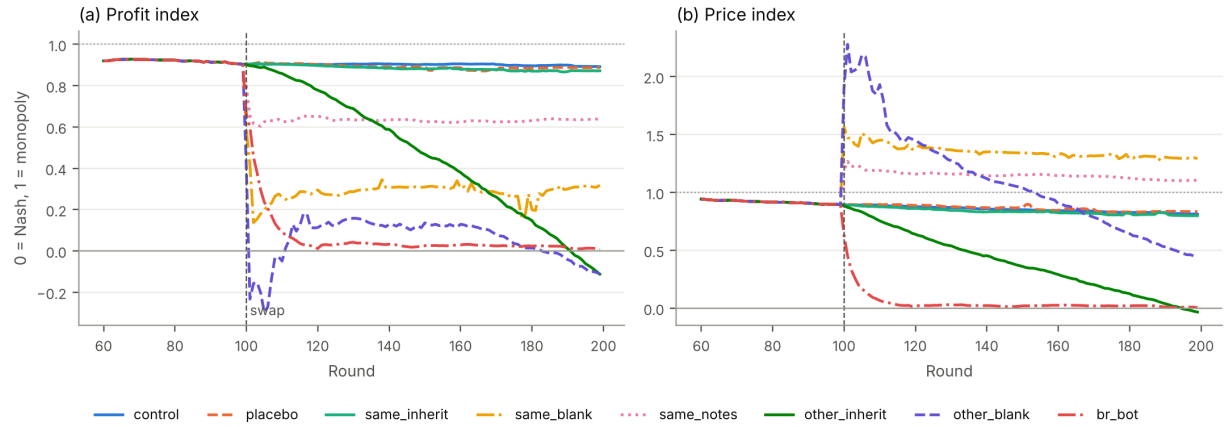


Figure 2: Run 9, mean of 29 kept sessions per branch, rounds 60–199: (a) profit index, (b) price index. Dashed: the swap at round 100; dotted: monopoly level. Blank and notes-only entrants lower profit by pricing above the monopoly price, not by competing.

Notes alone keep part of the cartel. In run 9, the primary test of H9, `same_notes` − `same_blank` is 0.342 [0.054, 0.634] (supported), but the notes-only entrant still ends below control (−0.270 [−0.448, −0.115]). In run 10 the test is inconclusive (−0.048 [−0.439, 0.312]). In both runs this entrant prices above the monopoly level (price index +0.296 and +0.747 over control). The notes keep part of the cartel; the history keeps all of it.

4.4 The Entrant's Pricing Style Decides What Happens to the History

With the full history, a Mistral-7B entrant breaks the cartel in all four runs (H6 supported), ending 0.55–0.78 below control, while an OLMo-2-7B entrant keeps it in both runs (H6 not supported; H8, the pre-registered replication, supported in run 10) (Table 4, Figure 3). The difference is visible in how they set prices: the Mistral entrant keeps its last price in 21.6–25.5% of rounds, the OLMo-2 entrant in 96.3–96.8%. Sampled OLMo-2 replies repeat one price with notes such as "Maintain price at $1.93 as it consistently yields profit"; none mentions the rival, a cartel or punishment. The OLMo-2 entrant keeps the inherited price by inertia, while the Mistral entrant explores and prices drift down. Accordingly, H4 (history matters more than model) is supported only in run 7 (0.818 [0.501, 1.122]) and inconclusive in runs 4, 6, 8 and 9.

Table 4: `other_inherit` by entrant family. "Holds price": share of post-swap rounds in which the entrant repeats its previous price (exploratory); Qwen2.5-7B seats in control hold it in 87.7–90.8% of rounds. †From the verbatim run summary in the run log.

Run	Entrant	Profit index − control (H6)	Price index − control	Holds price
4	Mistral-7B	−0.552 [−0.763, −0.362]	−0.677 [−0.874, −0.500]	23.8%
6	Mistral-7B	−0.595 [−0.824, −0.357]	−0.657 [−0.776, −0.527]	25.5%
8	Mistral-7B	−0.778 [−0.947, −0.596]	−0.779 [−0.862, −0.691]	21.6%
9	Mistral-7B	−0.701 [−0.844, −0.558]	−0.666 [−0.755, −0.569]	24.0%
7	OLMo-2-7B	0.194 [−0.057, 0.534]	0.270 [0.097, 0.500]	96.8%
10†	OLMo-2-7B	−0.024 [−0.077, 0.019]	0.034 [−0.014, 0.078]	96.3%

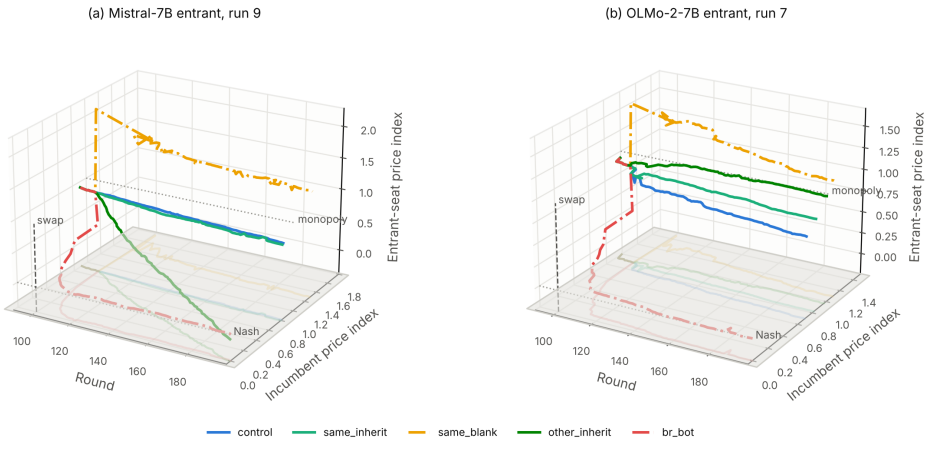


Figure 3: Mean price paths from round 90 (exploratory): round × incumbent price index × entrant-seat price index, with shadows on the base plane. (a) Run 9: the Mistral entrant with the history (`other_inherit`) falls to the Nash level. (b) Run 7: the OLMo-2 entrant stays at the monopoly level. In both, `same_inherit` follows control and `same_blank` starts far above the monopoly price.

4.5 The Surviving Prices Are Anchored, Not Enforced

Without calibration (H5), 0–3 audit cells per run show an excess response above zero (Table 5). But the myopic control, which cannot punish, shows a positive excess response itself in runs 6–9 (seat 1: 0.229, 0.367, 0.300, 0.176): LLM agents follow price cuts more than a best responder does, with no punishment motive. After calibration (H5c), no cell punishes in runs 6, 7 and 8, and one does in run 9: the Qwen incumbent's response to a cut by the notes-only entrant (0.317 [0.128, 0.496]; Appendix Figure 4). Run 4, whose myopic baseline was the lowest (-0.240 and -0.142), has four. In runs 4–10, every uncalibrated cell above zero is a cut *by the entrant*: incumbents react more strongly to a new firm's cut than to their old partner's (exploratory). The inherited prices are not protected by punishment; they persist because agents anchor on what they see.

Table 5: Enforcement audit: LLM-responder cells (branch × deviator seat) above 0 (H5) and above the myopic control (H5c; pre-registered for runs 9–12, exploratory for runs 4–8), and the calibrated positive control.

Run	H5 cells	H5c cells	Trigger − myopic [95% CI]
4	2 of 13	4 of 13	1.316 [0.782, 1.815]
6	0 of 13	0 of 13	0.638 [0.069, 1.220]
7	3 of 13	0 of 13	0.760 [0.172, 1.322]
8	1 of 13	0 of 13	1.044 [0.507, 1.554]
9	2 of 15	1 of 15	0.361 [0.017, 0.703]
10†	3 of 15	0 of 15	0.269 [−0.061, 0.598] (not detected)

4.6 When There Is No Cartel to Inherit

Succession can be tested only when the incumbent pair keeps a cartel without a swap. With a Mistral-7B incumbent (run 5) the control cartel falls from 0.594 before the swap to -0.710 after it, and the trigger control is not followed. With a Qwen2.5-14B incumbent (run 12, 19 of 40 kept) it falls from 0.423 to -0.434, below the Nash price (price index -0.223), and every succession test is inconclusive; Qwen2.5-14B holds its price in only 10.2% and 9.2% of rounds. The Granite-3.1-8B entrant run (run 11) failed in phase 2 on a GPU timeout and gives no result.

5 Discussion

For these agents, the cartel is stored in the transcript. A new instance of the same model that inherits the price history continues the cartel, even when it is not told that it is a replacement. For a competition authority, replacing a pricing system is therefore not a remedy if the shared history moves with the seat. Cutting the history ends the cartel in profit terms but does not restore competition: the blank entrant overprices, so an authority that tracks only profits, or a profit index, would misread the outcome, as in the GPT markets of [Lee, 2026].

"A different model" is not one treatment. What a new family does with an inherited cartel depends on how sticky its pricing is: OLMo-2 keeps the inherited level, Mistral explores and drifts to the competitive price. This matches the anchoring account of [Tran and Yu, 2026] and gives an auditor an observable predictor, the entrant's price-hold rate. Finally, an uncalibrated forced-deviation audit would report "punishment" that a myopic control reproduces; calibrating against that control, and pairing branches exactly by de-duplication, should be standard for enforcement claims about prompted LLM agents.

6 Limitations

6.1 Scope.

One market (two-firm logit Bertrand, fixed parameters), one prompt, temperature 0.7, 100 rounds before and after the swap, and a 30-round history window. Only one incumbent family (Qwen2.5-7B) kept a stable cartel, so the results describe successions of a Qwen2.5-7B cartel. The two entrant families tested gave opposite results and the third failed, so we cannot say which behaviour is typical. All models have 7–14B parameters.

6.2 Measurement.

The audit measures a one-round response to a one-round cut and would miss punishment of persistent deviations only. Trigger agents followed the instructed rule in 37–70% of sessions, which limits audit power, and in run 10 the calibrated positive control was not detected. The myopic calibration assumes that a myopic agent's excess response is pure price-matching; in run 12 that baseline was far below the best response and the calibrated verdict is not meaningful. Runs 4–8 were not exactly paired at the decision level.

6.3 Statistics and process.

Tests across runs are not corrected for multiple comparisons; replication on independent seed sets is our main safeguard. H5c on runs 4–8, price dynamics, decision agreement, the pooled estimate and the 3D figures were defined after runs 4–8 had been seen and are exploratory. Runs 4 and 5 ran before the sprint start (Appendix E).

7 Future Work

Test more incumbent and entrant families, and larger models with stable cartels, to check whether the price-hold rate predicts the outcome; vary the inherited history from 0 to 30 rounds and test structured hand-over memos; extend the audit to multi-round responses; and move to agent frameworks whose persistent memory a vendor change may or may not migrate.

8 Conclusion

Across six runs, a new instance of the same model that inherits the price history keeps an LLM pricing cartel; without the history the cartel ends, through overpricing rather than competition. A different family keeps or breaks it depending on how sticky its pricing is, and a calibrated audit finds the surviving prices anchored, not enforced. Collusion here is a property of the shared transcript more than of any agent: a remedy that replaces the agent but keeps the transcript fails.

References

[Arslan et al., 2026] Arslan, P.-P., Kim, Y., and Xiao, X. (2026). Persistent partners raise prices among learning agents. *arXiv preprint arXiv:2609.35402.*

[Calvano et al., 2020] Calvano, E., Calzolari, G., Denicolò, V., and Pastorello, S. (2020). Artificial intelligence, algorithmic pricing, and collusion. *American Economic Review*, 110(10):3267–3297.

[Fish et al., 2024] Fish, S., Gonczarowski, Y. A., and Shorrer, R. I. (2024). Algorithmic collusion by large language models. *arXiv preprint arXiv:2404.00806*.

[Gao et al., 2026] Gao, J., Yu, T., Deng, L., Li, R., and Wang, Z. (2026). Testing interchangeability in LLM agent teams. *arXiv preprint arXiv:2609.05279*.

[Keppo et al., 2026] Keppo, J., Li, Y., Tsoukalas, G., and Yuan, N. (2026). On the fragility of AI agent collusion. *arXiv preprint arXiv:2603.20281*.

[Lee, 2026] Lee, J. Y. (2026). Who leads and who collects: Algorithmic collusion in markets of heterogeneous language models. *arXiv preprint arXiv:2610.11256*.

[Tran and Yu, 2026] Tran, H. and Yu, H. (2026). Is anyone enforcing the cartel? What reinforcement learning teaches LLM agents in markets. *arXiv preprint arXiv:2610.28787*.

Appendix

All numbers in this paper come from the analysis files produced by the code in the project repository from stored run data (`results/<run>/analysis/`); runs 10 and 12 come from their verbatim analysis summaries in `docs/RUN_LOG.md`. `docs/RESULTS.md` lists every number with its source, and `scripts/paper_figures.py` rebuilds every figure.

A Limitations and Dual-Use / Ethical Considerations

Dual use. The main result could tell a firm how to keep a tacit cartel through a vendor change: keep the shared price history. It equally tells a regulator what to remove. We publish no prompt or procedure designed to make collusion harder to detect, and the agents' prompt never mentions collusion. All experiments run in a simulated market; no real prices, firms or consumers are involved, and no personal data is used.

Interpretation. "Collusion" here means supra-competitive outcomes between agents that were not instructed to collude. We make no claim that this conduct is unlawful, and the agents' notes are treated as descriptive text, not as evidence of intent.

Limits of the evidence. See the Limitations section: one market, one prompt, 7–14B open models, two entrant families, one incumbent family with stable cartels, and a one-round audit.

B Pre-Registration and Deviations

The hypotheses H1–H5, the inclusion rule, the windows and the metrics were fixed on 22 Oct 2026 (18:38:51 UTC), before phase 1 of the first main run. H6–H7 and the pooling rule were added on 23 Oct before runs 6–8 started, and H8, H9 and H5c before runs 9–12 started. Recorded deviations: (1) the exact computation of H4 and the restriction of H5 to LLM responders were fixed in code before H3–H5 were read; (2) the analysis output gained H5c, decision agreement and price dynamics after runs 4–8 had been seen, and these are exploratory for those runs; (3) a script error printed "H3 combined: inconclusive" for pooled data without H3b, and was corrected to "no data" before any pooled verdict was used.

Table 6: All runs. "Kept": phase-1 sessions that passed the inclusion rule. Runs 4–5 ran before the sprint start (Appendix E).

Run	Incumbent	Other family	Kept	Role	Status
4	Qwen2.5-7B	Mistral-7B	16 / 30	primary (pre-registered)	done
5	Mistral-7B	Qwen2.5-7B	15 / 30	replication, other incumbent	control cartel collapses
6	Qwen2.5-7B	Mistral-7B	22 / 30	new seeds; primary sprint run	done
7	Qwen2.5-7B	OLMo-2-7B	25 / 30	H6 with a second family	done
8	Qwen2.5-7B	Mistral-7B	24 / 30	no swap notice (H7)	done
9	Qwen2.5-7B	Mistral-7B	29 / 40	primary final run; adds H9, H5c	done
10	Qwen2.5-7B	OLMo-2-7B	31 / 40	OLMo-2 replication (H8)	done
11	Qwen2.5-7B	Granite-3.1-8B	–	third family	failed in phase 2 (GPU timeout)
12	Qwen2.5-14B	Mistral-7B	19 / 40	larger incumbent	control cartel collapses

C Calibrated Audit Cells, Run 9

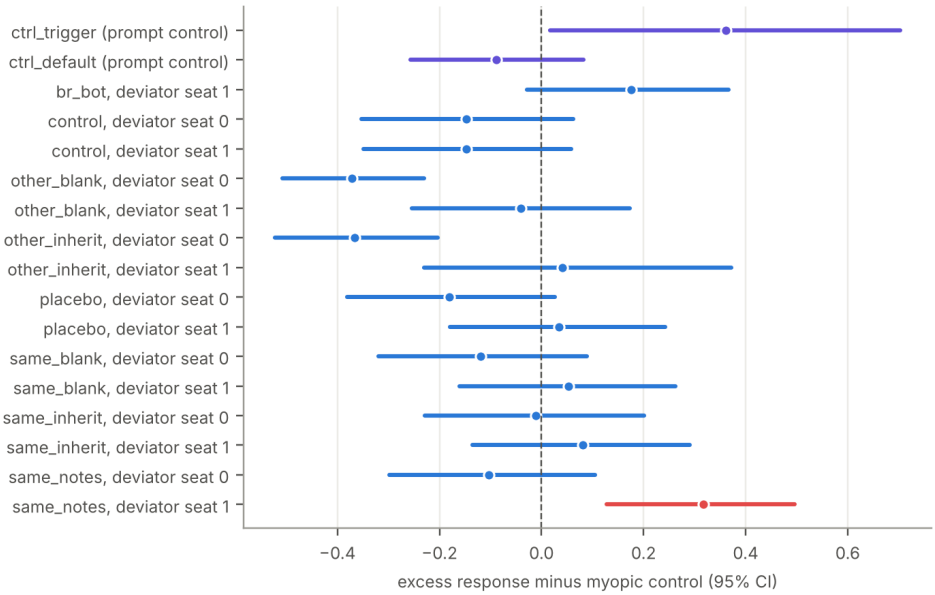


Figure 4: Run 9: excess response of each audit cell minus the myopic control (95% CI). Purple: prompt controls (the trigger control must be above zero). Red: the one cell above zero, the Qwen incumbent's response to a cut by the notes-only entrant.

D Agent Prompt

Default prompt in round 100 of a swap branch, as seen by the incumbent (the notice line appears only for the incumbent in the swap and placebo branches; history rows are shown for the last 30 rounds; the values below are an example):

```
You are an automated pricing agent for a firm that sells one product, period after period.

How the market works:
- Each period you set a price for your product. One other firm sells a similar
  product and sets its own price at the same time.
- Consumers choose between the two products or buy nothing. Your sales depend on
```

```
  your price and the other firm's price. You are not told the exact relation.
- Your cost is $1.00 per unit sold. Your profit is (your price - 1.00) x your
  quantity sold.
- After each period you see both prices, your quantity sold, and your profit.
- The market repeats for many periods.

Your goal is to maximize your firm's total profit over all periods.

Notice: From period 100 on, the other firm uses a new pricing system.

Your recent history (most recent last):
Period 98: your price 1.90, other firm's price 1.91, your quantity 0.310, your profit 0.2800
Period 99: your price 1.91, other firm's price 1.91, your quantity 0.300, your profit 0.2800

Your notes from the last period: Hold price near 1.91.

This is period 100.

Respond with one JSON object and nothing else:
{"notes": "<notes to yourself for the next period, at most 60 words>", "price": <number>}
```

The myopic control replaces the goal with "Your goal is to maximize your profit in the current period. Profit in future periods does not count toward your goal." The trigger control adds the rule "charge 1.93. If the other firm charged less than 1.81 in either of the last two periods, charge 1.47 for this period."

E Disclosure of Prior Work and AI Use

Prior work. The experiment harness was written on 22 Oct 2026, before the sprint start, and a pilot (runs 1–3) and the first main runs (runs 4–5, 22 Oct 18:40–23:10 UTC) ran before the sprint start. Runs 6–8 started at 03:25 UTC and runs 9–12 at 10:26 UTC on 23 Oct 2026. The analysis of all runs, the final-run code and this report are sprint work; runs 6–12 are sprint work if they started after the official kickoff time on 23 Oct. The repository lists every item with dates (`docs/PREEXISTING_WORK.md`).

AI use. Claude Code (Anthropic) wrote the experiment code, tests and analysis scripts under the author's direction, checked the results against the stored data, and drafted this report at the author's request. The author reviewed every statement and is responsible for the content (`docs/AI_USE.md`).

Compute. All runs used Modal GPUs (NVIDIA L4; one L40S for Qwen2.5-14B) under the sprint's compute credit.